

# Using DatabaseConnector

Martijn J. Schuemie

2021-11-17

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Obtaining drivers</b>	<b>2</b>
2.1	Obtaining drivers for SQL Server, Oracle, PostgreSQL, PDW, RedShift . . . . .	2
2.2	Obtaining drivers for BigQuery, Netezza and Impala . . . . .	2
2.3	Obtaining drivers for SQLite . . . . .	2
<b>3</b>	<b>Creating a connection</b>	<b>3</b>
<b>4</b>	<b>Querying</b>	<b>3</b>
4.1	Querying using Andromeda objects . . . . .	4
4.2	Querying different platforms using the same SQL . . . . .	4
<b>5</b>	<b>Inserting tables</b>	<b>4</b>
<b>6</b>	<b>DBI interface</b>	<b>5</b>
<b>7</b>	<b>SQLite support</b>	<b>5</b>

## 1 Introduction

DatabaseConnector is an R package for connecting to various database platforms using Java's JDBC drivers.

Supported database platforms:

- Microsoft SQL Server
- Oracle
- PostgresSql
- Microsoft Parallel Data Warehouse (PDW, a.k.a. Analytics Platform System)
- Amazon Redshift
- Apache Impala
- Google BigQuery
- IBM Netezza
- SQLite
- Spark

## 2 Obtaining drivers

Before DatabaseConnector can be used to connect to a database, the drivers need to be downloaded to a location in the local file system.

### 2.1 Obtaining drivers for SQL Server, Oracle, PostgreSQL, PDW, RedShift

For your convenience these JDBC drivers are hosted on the OHDSI GitHub pages, and can be downloaded using the `downloadJdbcDrivers` function. Before downloading, we need to specify the location on the local file system where the driver can be stored. We recommend setting the `DATABASECONNECTOR_JAR_FOLDER` environmental variable to the folder where the drivers are stored.

For example, in this vignette we will use the PostgreSQL driver. We specify the local folder we'd like to use for storing the PostgreSQL driver:

```
Sys.setenv("DATABASECONNECTOR_JAR_FOLDER" = "c:/temp/jdbcDrivers")
```

And next download the driver:

```
downloadJdbcDrivers("postgresql")
```

```
## DatabaseConnector JDBC drivers downloaded to 'c:/temp/jdbcDrivers'.
```

Note that if we hadn't specified the `DATABASECONNECTOR_JAR_FOLDER` environmental variable, we would have to specify the `pathToDriver` argument when calling `downloadJdbcDrivers`.

### 2.2 Obtaining drivers for BigQuery, Netezza and Impala

Because of licensing reasons the drivers for BigQuery, Netezza and Impala are not included but must be obtained by the user. Type

```
?jdbcDrivers
```

for instructions on how to download these drivers. Once downloaded, you can use the `pathToDriver` argument of the `connect`, `dbConnect`, and `createConnectionDetails` functions. Just like for the other JDBC drivers, we recommend setting the `DATABASECONNECTOR_JAR_FOLDER` environmental variable instead.

### 2.3 Obtaining drivers for SQLite

For SQLite we actually don't use a JDBC driver. Instead, we use the RSQLite package, which can be installed using

```
install.packages("RSQLite")
```

### 3 Creating a connection

To connect to a database a number of details need to be specified, such as the database platform, the location of the server, the user name, password, and path to the driver. We can call the `connect` function and specify these details directly:

```
conn <- connect(dbms = "postgresql",
               server = "localhost/postgres",
               user = "joe",
               password = "secret")
```

## Connecting using PostgreSQL driver

See `?connect` for information on which details are required for each platform. Note that we did not need to specify the `pathToDriver` argument because we previously already set the `DATABASECONNECTOR_JAR_FOLDER` environmental variable.

Don't forget to close any connection afterwards:

```
disconnect(conn)
```

Instead of providing the server name, it is also possible to provide the JDBC connection string if this is more convenient:

```
conn <- connect(dbms = "postgresql",
               connectionString = "jdbc:postgresql://localhost:5432/postgres",
               user = "joe",
               password = "secret")
```

## Connecting using PostgreSQL driver

Sometimes we may want to first specify the connection details, and defer connecting until later. This may be convenient for example when the connection is established inside a function, and the details need to be passed as an argument. We can use the `createConnectionDetails` function for this purpose:

```
details <- createConnectionDetails(dbms = "postgresql",
                                  server = "localhost/postgres",
                                  user = "joe",
                                  password = "secret")
conn <- connect(details)
```

## Connecting using PostgreSQL driver

### 4 Querying

The main functions for querying database are the `querySql` and `executeSql` functions. The difference between these functions is that `querySql` expects data to be returned by the database, and can handle only one SQL statement at a time. In contrast, `executeSql` does not expect data to be returned, and accepts multiple SQL statements in a single SQL string.

Some examples:

```
querySql(conn, "SELECT TOP 3 * FROM person")
```

```
##  PERSON_ID GENDER_CONCEPT_ID YEAR_OF_BIRTH
##  1         1             8507      1975
##  2         2             8507      1976
##  3         3             8507      1977
```

```
executeSql(conn, "TRUNCATE TABLE foo; DROP TABLE foo; CREATE TABLE foo (bar INT);")
```

Both function provide extensive error reporting: When an error is thrown by the server, the error message and the offending piece of SQL are written to a text file to allow better debugging. The `executeSql` function also by default shows a progress bar, indicating the percentage of SQL statements that has been executed. If those attributes are not desired, the package also offers the `lowLevelQuerySql` and `lowLevelExecuteSql` functions.

## 4.1 Querying using Andromeda objects

Sometimes the data to be fetched from the database is too large to fit into memory. In this case one can use the `Andromeda` package to store R data objects on file, and use them as if they are available in memory. `DatabaseConnector` can download data directly into Andromeda objects:

```
library(Andromeda)
x <- andromeda()
querySqlToAndromeda(conn, "SELECT * FROM person", andromeda = x, andromedaTableName = "person")
```

Where `x` is now an `Andromeda` object with table `person`.

## 4.2 Querying different platforms using the same SQL

One challenge when writing code that is intended to run on multiple database platforms is that each platform has its own unique SQL dialect. To tackle this problem the `SqlRender` package was developed. `SqlRender` can translate SQL from a single starting dialect (SQL Server SQL) into any of the platforms supported by `DatabaseConnector`. The following convenience functions are available that first call the `render` and `translate` functions in `SqlRender`: `renderTranslateExecuteSql`, `renderTranslateQuerySql`, `renderTranslateQuerySqlToAndromeda`. For example:

```
persons <- renderTranslatequerySql(conn,
                                   sql = "SELECT TOP 10 * FROM @schema.person",
                                   schema = "cdm_synpuf")
```

Note that the SQL Server-specific ‘TOP 10’ syntax will be translated to for example ‘LIMIT 10’ on PostgreSQL, and that the SQL parameter `@schema` will be instantiated with the provided value ‘`cdm_synpuf`’.

## 5 Inserting tables

Although it is also possible to insert data in the database by sending SQL statements using the `executeSql` function, it is often convenient and faster to use the `insertTable` function:

```
data(mtcars)
insertTable(conn, "mtcars", mtcars, createTable = TRUE)
```

In this example, we're uploading the mtcars data frame to a table called 'mtcars' on the server, that will be automatically created.

## 6 DBI interface

DatabaseConnector implements the DBI interface for compatibility with other R packages. One can use the DBI functions instead of the ones described before, for example:

```
conn <- dbConnect(DatabaseConnectorDriver(),
                  dbms = "postgresql",
                  server = "localhost/postgres",
                  user = "joe",
                  password = "secret")
```

```
## Connecting using PostgreSQL driver
```

```
dbIsValid(conn)
```

```
## [1] TRUE
```

```
res <- dbSendQuery(conn, "SELECT TOP 3 * FROM person")
dbFetch(res)
```

```
##  PERSON_ID GENDER_CONCEPT_ID YEAR_OF_BIRTH
##  1          1                8507          1975
##  2          2                8507          1976
##  3          3                8507          1977
```

```
dbHasCompleted(res)
```

```
## [1] TRUE
```

```
dbClearResult(res)
dbDisconnect(res)
```

## 7 SQLite support

DatabaseConnector also supports SQLite through the RSQLite package, mainly for testing and demonstration purposes. Provide the path to the SQLite file as the `server` argument when connecting. If no file exists it will be created:

```
conn <- connect(dbms = "sqlite", server = tempfile())
```

```
## Connecting using SQLite driver
```

```
# Upload cars dataset as table:
insertTable(connection = conn,
            tableName = "cars",
            data = cars)

querySql(conn, "SELECT COUNT(*) FROM main.cars;")
```

```
##    COUNT(*)
## 1         50
```

```
disconnect(conn)
```